

✓ HW4 Oliver Li

```
1 ## Problem 1 - Recent Posts
2
3 <change you made>
4
5 CREATE INDEX post_timestamp_idx ON posts(post_timestamp);
6 SELECT
7     post_id,
8     post_timestamp
9
10 FROM
11     posts
12
13 ORDER BY
14     post_timestamp DESC
15
16 LIMIT 10
17
18 <screenshot of EXPLAIN ANALYZE>
```

PLAN

```
SELECT
  "POST_ID",
  "POST_TIMESTAMP"
FROM "PUBLIC"."POSTS"
/* PUBLIC.POST_TIMESTAMP_IDX */
/* scanCount: 10 */
ORDER BY 2 DESC
FETCH FIRST 10 ROWS ONLY
/* index sorted */
/*
reads: 4
*/
(1 row, 7 ms)
```

Adding an index to improve performance

```
1 ## Problem 2 - Somewhat Strange Query
2
3 <change you made>
4 ALTER TABLE posts
5     ADD COLUMN author_substr VARCHAR(3) AS SUBSTR(author, 3, 3);
6
7 ALTER TABLE posts
8     ADD COLUMN content_upper VARCHAR(255) AS UPPER(content);
9
10 CREATE INDEX content_idx ON posts (content);
11 CREATE INDEX post_timestamp_idx ON posts (post_timestamp);
12 CREATE INDEX author_substr_idx ON posts (author_substr);
13 CREATE INDEX content_upper_idx ON posts (content_upper);
14 SELECT
15     post_id,
16     post_timestamp
17 FROM
18     posts
19 WHERE
20     post_timestamp < '2024-02-01'
21     AND content_upper LIKE 'C%'
22     AND author_substr = 'son';
23
24 <screenshot of EXPLAIN ANALYZE>
```

```

PLAN
SELECT
  "POST_ID",
  "POST_TIMESTAMP"
FROM "PUBLIC"."POSTS"
/* PUBLIC.POSTS:tableScan */
/* scanCount: 995087 */
WHERE ("AUTHOR_SUBSTR" = 'son')
AND ("POST_TIMESTAMP" < TIMESTAMP '2024-02-01 00:00:00')
AND ("CONTENT_UPPER" LIKE 'C%')
/*
reads: 101444
*/
(1 row, 768 ms)

```

EXPLAIN ANALYZE

```

SELECT
  post_id,
  post_timestamp

FROM
  posts

WHERE
  post_timestamp < '2024-02-01'
  AND content_upper LIKE 'C%'
  AND author_substr = 'son';

```

PLAN

```

SELECT
  "POST_ID",
  "POST_TIMESTAMP"
FROM "PUBLIC"."POSTS"
/* PUBLIC.AUTHOR_SUBSTR_IDX: AUTHOR_SUBSTR = 'son' */
/* scanCount: 35586 */
WHERE ("AUTHOR_SUBSTR" = 'son')
AND ("POST_TIMESTAMP" < TIMESTAMP '2024-02-01 00:00:00')
AND ("CONTENT_UPPER" LIKE 'C%')
/*
reads: 5666
*/
(1 row, 62 ms)

```

Creating an index on all the provided queries allows faster performance speed. Also creating subrows to provide the indices.

✓ Problem 3 - Really Fast Single Row Responses

Problem 3.1

<What index does H2DB end up using? Explain the pros and cons of each index that you created.> The H2DB ends up using the B-Tree indexing. For both of the indices: B-Trees work fast on range queries and sorting, perform well ($O(\log N)$) in lookup tasks/exact match queries, and more general to use. However B-Trees take up more space and memory in practice, also performs slower in insets and updates due to re-balancing in data. On the other hand Hash Indices use less space and performs fastest on lookup time. However it is neither useful for range queries or sorting queries.

```

SET TRACE_LEVEL_SYSTEM_OUT 3;
SET CACHE_SIZE 0;
SET MAX_MEMORY_ROWS 0;
SET MAX_MEMORY_UNDO 0;
EXPLAIN ANALYZE

SELECT
  post_id,
  post_timestamp

FROM
  posts

WHERE
  post_timestamp = '2024-01-26 17:52:23.000000';

CREATE INDEX post_timestamp_idx ON posts (post_timestamp);

DROP INDEX post_timestamp_idx;

```

```

EXPLAIN ANALYZE

SELECT
  post_id,
  post_timestamp

FROM
  posts

WHERE
  post_timestamp = '2024-01-26 17:52:23.000000';
PLAN
SELECT
  "POST_ID",
  "POST_TIMESTAMP"
FROM "PUBLIC"."POSTS"
/* PUBLIC.IDX_POST_TIMESTAMP: POST_TIMESTAMP = TIMESTAMP '2024-01-26 17:52:23' */
/* scanCount: 4 */
WHERE "POST_TIMESTAMP" = TIMESTAMP '2024-01-26 17:52:23'
/*
reads: 4
*/
(1 row, 7 ms)

```

Problem 3.2

<Which of the indexes that you created for 3.1 would you expect to be used now. Please explain.> It is expected to use a B-Tree Index. A BETWEEN AND operation is not supported by a hash index, as hash indices do not have order.

```

SET TRACE_LEVEL_SYSTEM_OUT 3;
SET CACHE_SIZE 0;
SET MAX_MEMORY_ROWS 0;
SET MAX_MEMORY_UNDO 0;
EXPLAIN ANALYZE

SELECT
    post_id,
    post_timestamp

FROM
    posts

WHERE
    post_timestamp BETWEEN '2024-01-25' AND '2024-01-27';

CREATE HASH INDEX post_hash_timestamp_idx ON posts (post_timestamp);
CREATE INDEX post_btree_timestamp_idx ON posts (post_timestamp);

DROP INDEX post_btree_timestamp_idx ;
DROP INDEX post_hash_timestamp_idx ;

```

```

SELECT
    post_id,
    post_timestamp

FROM
    posts

WHERE
    post_timestamp BETWEEN '2024-01-25' AND '2024-01-27';
PLAN
SELECT
    "POST_ID",
    "POST_TIMESTAMP"
FROM "PUBLIC"."POSTS"
/* PUBLIC.POST_HASH_TIMESTAMP_IDX: POST_TIMESTAMP >= '2024-01-25'
   AND POST_TIMESTAMP <= '2024-01-27'
*/
/* scanCount: 1000 */
WHERE "POST_TIMESTAMP" BETWEEN '2024-01-25' AND '2024-01-27'
/*
reads: 45
*/
(1 row, 7 ms)

```

Problem 3.3

<Can you modify one of the indexes from 3.2 to make this query even faster? Explain why your change to the index made the query even faster.> CREATE INDEX post_timestamp_btree_idx ON posts (post_timestamp, post_id, content); This is a composite index that allow he coverage of all the columns in the query, where the database can fetch data directly from the index without having to access the full table.

```

1 ## Problem 4 - Table Join Order
2 ### Problem 4.1
3
4 <Your modified query here>
5
6 SELECT
7     COUNT(1)
8
9 FROM
10     users
11     JOIN followers ON users.handle = followers.follower_handle
12     JOIN posts ON followers.following_handle = posts.author
13
14 WHERE
15     users.last_name = 'Anderson'
16     AND users.first_name = 'Abigail';

```

Problem 4.2

We have 3 tables to join hence we have $3!-2=4$ possible joins possible, and known that the size of each table is: users = 10000; followers = 995040; posts = 995086; Note that the performance can only be optimized when each operation eliminates most elements first, hence given possible joins:

```

1 FROM
2     posts
3     JOIN followers ON posts.author = followers.following_handle
4     JOIN users ON followers.follower_handle = users.handle
5 ### Gives 205 in 103754 ms
6 ### Likely to take the longest time as the operations take the longest table first, then the second longest and the shortest.
7 ### Elimination of rows is minimized.
8
9 FROM
10     users
11     JOIN followers ON users.handle = followers.follower_handle
12     JOIN posts ON followers.following_handle = posts.author
13 ### Gives 205 in 10 ms
14 ### Likely to take the shortest time as the operations take the shortest table first, then the second shortest and the longest.
15 ### Elimination of rows is maximized as the table is short in the first place.
16
17 FROM
18     followers
19     JOIN posts ON followers.following_handle = posts.author
20     JOIN users ON followers.follower_handle = users.handle
21 ### Gives 205 in 61401 ms
22 ### Better than the first JOIN order, as the second longest table is given at first, then the longest and the shortest
23 ### Elimination of rows is not maximized.
24
25 FROM
26     followers
27     JOIN users ON followers.follower_handle = users.handle
28     JOIN posts ON followers.following_handle = posts.author
29 ### Gives 205 in 672 ms
30 ### Better than the previous JOIN order, as the second longest table is given at first, then the shortest and the longest
31 ### Elimination of rows is not maximized but slightly better than the previous operation order, as 995040->10000 is better than 99

```

```

1 ## Problem 5 - Putting it All Together - Fast Most Recent Posts
2
3 <your query here>
4 DROP INDEX IF EXISTS posts_author_idx;
5 DROP INDEX IF EXISTS followers_follower_idx;
6 DROP INDEX IF EXISTS followers_following_idx;
7 DROP INDEX IF EXISTS posts_timestamp_idx;
8
9 CREATE HASH INDEX followers_following_idx ON followers (following_handle);
10 CREATE HASH INDEX followers_follower_idx ON followers (follower_handle);
11 CREATE HASH INDEX posts_author_idx ON posts (author);
12
13 CREATE INDEX posts_timestamp_idx ON posts (post_timestamp DESC);
14
15 EXPLAIN ANALYZE
16 WITH latest_posts AS (
17     SELECT
18         pl.author,
19         MAX(pl.post_timestamp) AS latest
20     FROM followers f
21     JOIN posts pl
22         ON pl.author = f.following_handle
23     WHERE f.follower_handle = 'madison.anderson9901'
24     GROUP BY pl.author
25 )
26 SELECT
27     p2.author,
28     p2.post_id,
29     p2.post_timestamp,
30     p2.content
31 FROM latest_posts sub
32 JOIN posts p2
33     ON p2.author = sub.author
34     AND p2.post_timestamp = sub.latest
35 ORDER BY p2.post_timestamp DESC;
36

```

PLAN

```
WITH "LATEST_POSTS"("AUTHOR", "LATEST") AS (  
  SELECT  
    "P1"."AUTHOR",  
    MAX("P1"."POST_TIMESTAMP") AS "LATEST"  
  FROM "PUBLIC"."FOLLOWERS" "F"  
  INNER JOIN "PUBLIC"."POSTS" "P1"  
    ON 1=1  
  WHERE ("F"."FOLLOWER_HANDLE" = 'madison.anderson9901')  
    AND ("P1"."AUTHOR" = "F"."FOLLOWING_HANDLE")  
  GROUP BY "P1"."AUTHOR"  
)  
SELECT  
  "P2"."AUTHOR",  
  "P2"."POST_ID",  
  "P2"."POST_TIMESTAMP",  
  "P2"."CONTENT"  
FROM "LATEST_POSTS" "SUB"  
/* SELECT  
  P1.AUTHOR,  
  MAX(P1.POST_TIMESTAMP) AS LATEST  
FROM PUBLIC.FOLLOWERS F  
/* PUBLIC.PRIMARY_KEY_D: FOLLOWER_HANDLE = 'madison.anderson9901' */  
/* WHERE F.FOLLOWER_HANDLE = 'madison.anderson9901'  
*/  
/* scanCount: 5880 */  
INNER JOIN PUBLIC.POSTS P1  
/* PUBLIC.CONSTRAINT_48C_INDEX_D: AUTHOR = F.FOLLOWING_HANDLE */  
ON 1=1  
/* scanCount: 577000 */  
WHERE (F.FOLLOWER_HANDLE = 'madison.anderson9901')  
  AND (P1.AUTHOR = F.FOLLOWING_HANDLE)  
GROUP BY P1.AUTHOR  
*/  
/* scanCount: 5752 */  
INNER JOIN "PUBLIC"."POSTS" "P2"  
/* PUBLIC.POSTS_TIMESTAMP_IDX: POST_TIMESTAMP = SUB.LATEST */  
ON 1=1  
/* scanCount: 11686 */
```